

Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import csv
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import MinMaxScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from scipy.stats import pearsonr
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error
import statsmodels.api as sm
from xgboost import XGBRegressor
from sklearn.model_selection import GridSearchCV
```

```
In [2]: import sys
import os
sys.path.append(os.path.dirname(os.path.dirname(os.getcwd())))
if not os.path.exists("./data"):
    os.makedirs("./data")
if not os.path.exists("./data/model"):
    os.makedirs("./data/model")
if not os.path.exists("./data/preprocessing"):
    os.makedirs("./data/preprocessing")
if not os.path.exists("./data/dataset"):
    os.makedirs("./data/dataset")
```

Reusable Function

```
In [21]: def boxplot_los_groupby(variable, los_range=(-1, 30), size=(8,4)):
    results = df[[variable, 'los']].groupby(variable).median().reset_index()

    categories = results[variable].values.tolist()

    hist_data = []
    for cat in categories:
        hist_data.append(df['los'].loc[df[variable]==cat].values)

    fig, ax = plt.subplots(figsize=size)
    ax.boxplot(hist_data, 0, '', vert=False)
    ax.set_xlim(los_range)
    ax.set_yticklabels(categories)
    ax.set_xlabel('Length of Stay (days)')
    ax.tick_params(left=False, right=False)
    ax.set_title('Comparison of {} categories'.format(variable))
    plt.tight_layout()
    plt.show()
```

Extract Patients Data from ICUSTays

```
In [3]: icu = pd.read_csv('./mimiciv/2.0/icu/icustays.csv.gz',compression='gzip')
adm = pd.read_csv('./mimiciv/2.0/hosp/admissions.csv.gz',compression='gzip')
patients = pd.read_csv('./mimiciv/2.0/hosp/patients.csv.gz',compression='gzip')
```

```
In [4]: icu.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 73181 entries, 0 to 73180
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   subject_id      73181 non-null  int64
1   hadm_id         73181 non-null  int64
2   stay_id         73181 non-null  int64
3   first_careunit  73181 non-null  object
4   last_careunit   73181 non-null  object
5   intime          73181 non-null  object
6   outtime         73181 non-null  object
7   los             73181 non-null  float64
dtypes: float64(1), int64(3), object(4)
memory usage: 4.5+ MB
```

In [5]: `adm.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 431231 entries, 0 to 431230
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject_id            431231 non-null  int64
1   hadm_id               431231 non-null  int64
2   admittime             431231 non-null  object
3   dischtime            431231 non-null  object
4   deathtime            8598 non-null    object
5   admission_type        431231 non-null  object
6   admit_provider_id     431227 non-null  object
7   admission_location    431231 non-null  object
8   discharge_location    312076 non-null  object
9   insurance             431231 non-null  object
10  language              431231 non-null  object
11  marital_status        421998 non-null  object
12  race                  431231 non-null  object
13  edregtime             299282 non-null  object
14  edouttime             299282 non-null  object
15  hospital_expire_flag  431231 non-null  int64
dtypes: int64(3), object(13)
memory usage: 52.6+ MB
```

In [6]: *#drop unnecessary columns*

```
drop_adm = adm.drop(columns=['admittime', 'dischtime', 'deathtime', 'admission_type',
                             'admit_provider_id', 'admission_location', 'discharge_location',
                             'insurance', 'language', 'marital_status', 'race', 'edregtime', 'edouttime'])
```

In [7]: `drop_adm.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 431231 entries, 0 to 431230
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject_id            431231 non-null  int64
1   hadm_id               431231 non-null  int64
2   hospital_expire_flag  431231 non-null  int64
dtypes: int64(3)
memory usage: 9.9 MB
```

In [8]: *#merge icu with filtered admission*

```
df = icu.merge(drop_adm[['hadm_id', 'hospital_expire_flag']], on='hadm_id', how='left')
```

In [9]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 73181 entries, 0 to 73180
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   subject_id            73181 non-null  int64
 1   hadm_id               73181 non-null  int64
 2   stay_id              73181 non-null  int64
 3   first_careunit        73181 non-null  object
 4   last_careunit         73181 non-null  object
 5   intime               73181 non-null  object
 6   outtime              73181 non-null  object
 7   los                  73181 non-null  float64
 8   hospital_expire_flag  73181 non-null  int64
dtypes: float64(1), int64(4), object(4)
memory usage: 5.6+ MB
```

In [10]: `df.head()`

Out[10]:

	subject_id	hadm_id	stay_id	first_careunit	last_careunit	intime	outtime	los
0	10000032	29079034	39553978	Medical Intensive Care Unit (MICU)	Medical Intensive Care Unit (MICU)	2180- 07-23 14:00:00	2180- 07-23 23:50:47	0.410266
1	10000980	26913865	39765666	Medical Intensive Care Unit (MICU)	Medical Intensive Care Unit (MICU)	2189- 06-27 08:42:00	2189- 06-27 20:38:27	0.497535
2	10001217	24597018	37067082	Surgical Intensive Care Unit (SICU)	Surgical Intensive Care Unit (SICU)	2157- 11-20 19:18:02	2157- 11-21 22:08:00	1.118032
3	10001217	27703517	34592300	Surgical Intensive Care Unit (SICU)	Surgical Intensive Care Unit (SICU)	2157- 12-19 15:42:24	2157- 12-20 14:27:41	0.948113
4	10001725	25563031	31205490	Medical/Surgical Intensive Care Unit (MICU/SICU)	Medical/Surgical Intensive Care Unit (MICU/SICU)	2110- 04-11 15:52:22	2110- 04-12 23:59:56	1.338588

In [11]: `#Filter the 'icu' DataFrame to include only patients with at Least 5 hours of`
`filtered_icu = icu[icu['los'] > 5/24]`

In [12]: `#Merge the 'patients' and filtered 'icu' DataFrames based on the 'subject_id'`
`filtered_pats = patients.merge(filtered_icu[['subject_id',`
 `'intime', 'los']],on='subject_id'`

In [13]: `filtered_pats.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 72390 entries, 0 to 72389
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject_id            72390 non-null  int64
1   gender                 72390 non-null  object
2   anchor_age             72390 non-null  int64
3   anchor_year            72390 non-null  int64
4   anchor_year_group      72390 non-null  object
5   dod                    27908 non-null  object
6   intime                 72390 non-null  object
7   los                    72390 non-null  float64
dtypes: float64(1), int64(3), object(4)
memory usage: 5.0+ MB
```

In [14]: `# Filter out anyone whose age is less than or equal to 17 at the time of ICU admission`
`filtered_pats = filtered_pats[filtered_pats['anchor_age'] > 17]`

In [15]: `filtered_pats.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 72390 entries, 0 to 72389
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject_id            72390 non-null  int64
1   gender                 72390 non-null  object
2   anchor_age             72390 non-null  int64
3   anchor_year            72390 non-null  int64
4   anchor_year_group      72390 non-null  object
5   dod                    27908 non-null  object
6   intime                 72390 non-null  object
7   los                    72390 non-null  float64
dtypes: float64(1), int64(3), object(4)
memory usage: 5.0+ MB
```

In [16]: `df = df.merge(filtered_pats[['subject_id', 'gender', 'anchor_age']], on='subject_id', how='left')`


In [18]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 170453 entries, 0 to 170452
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   subject_id            170453 non-null  int64
 1   hadm_id                170453 non-null  int64
 2   stay_id                170453 non-null  int64
 3   first_careunit         170453 non-null  object
 4   last_careunit          170453 non-null  object
 5   intime                 170453 non-null  object
 6   outtime                170453 non-null  object
 7   los                    170453 non-null  float64
 8   hospital_expire_flag   170453 non-null  int64
 9   gender                 170453 non-null  object
10   anchor_age             170453 non-null  int64
dtypes: float64(1), int64(5), object(5)
memory usage: 15.6+ MB
```

In [17]: `#save icu cohort csv file`
`df.to_csv('./data/preprocessing/icu_cohort.csv.gz', index=False)`

```
In [19]: #plot for patients age vs los
plt.scatter(df['anchor_age'], df['los'], alpha=0.1)
plt.ylabel('LOS (days)')
plt.xlabel('Age (years)')
plt.title('Age versus Length-of-stay')
plt.ylim(1, 120)
```


Out[19]: (1.0, 120.0)



Scatter Plot Observation

- The average LOS for patients in their 60s is around 10 days, but there are some patients in this age group who have an LOS of less than 5 days or more than 20 days.
- The median LOS is around 10 days.
- There are a few outliers (Higher than average LOS) in the data.

```
In [22]: boxplot_los_groupby('first_careunit', los_range=(-5,30))
```



Labs Stay

extracting the itemids for all the labevents that occur within the time bounds for our cohort

```
In [3]: icu_cohort = pd.read_csv('./data/preprocessing/icu_cohort.csv.gz',compression=
```

```
In [4]: labitems = pd.read_csv('./mimiciv/2.0/hosp/d_labitems.csv.gz',compression='gzi
```

```
In [27]: labitems.head()
```

```
Out[27]:
```

	itemid	label	fluid	category
0	50801	Alveolar-arterial Gradient	Blood	Blood Gas
1	50802	Base Excess	Blood	Blood Gas
2	50803	Calculated Bicarbonate, Whole Blood	Blood	Blood Gas
3	50804	Calculated Total CO2	Blood	Blood Gas
4	50805	Carboxyhemoglobin	Blood	Blood Gas

```
In [28]: category_counts = labitems.groupby('category')['itemid'].count()

print(category_counts)
```

```
category
Blood Gas      64
Chemistry     777
Hematology     781
Name: itemid, dtype: int64
```



```

In [29]: file_path = './mimiciv/2.0/hosp/labevents.csv.gz'

chunk_size = 100000
chunks = []

for chunk in pd.read_csv(file_path, chunksize=chunk_size):
    merge_chunk = chunk.merge(icu_cohort[['stay_id', 'hadm_id', 'intime', 'los

    # Filter only numerical data
    merge_chunk = merge_chunk[merge_chunk['valuenum'].notnull()]

    # Convert the 'charttime' and 'intime' columns to pandas DateTime objects
    merge_chunk['charttime'] = pd.to_datetime(merge_chunk['charttime'])
    merge_chunk['intime'] = pd.to_datetime(merge_chunk['intime'])

    # Calculate the time difference between 'charttime' and 'intime' in days
    # we want to extract measurements between admission and the end of the pat
    merge_chunk['time_diff_days'] = (merge_chunk['charttime'] - merge_chunk['i

    # Filter the rows in 'merge_chunk'
    merge_chunk = merge_chunk[(merge_chunk['time_diff_days'] >= -1) & (merge_c

    chunks.append(merge_chunk)

labsstay = pd.concat(chunks, ignore_index=True)

```

C:\Users\User\AppData\Local\Temp\ipykernel_7652\1204666657.py:6: DtypeWarni
ng: Columns (5) have mixed types. Specify dtype option on import or set low
_memory=False.

```
for chunk in pd.read_csv(file_path, chunksize=chunk_size):
```

C:\Users\User\AppData\Local\Temp\ipykernel_7652\1204666657.py:6: DtypeWarni
ng: Columns (5) have mixed types. Specify dtype option on import or set low
_memory=False.

```
for chunk in pd.read_csv(file_path, chunksize=chunk_size):
```

C:\Users\User\AppData\Local\Temp\ipykernel_7652\1204666657.py:6: DtypeWarni
ng: Columns (5) have mixed types. Specify dtype option on import or set low
_memory=False.

```
for chunk in pd.read_csv(file_path, chunksize=chunk_size):
```

C:\Users\User\AppData\Local\Temp\ipykernel_7652\1204666657.py:6: DtypeWarni
ng: Columns (5) have mixed types. Specify dtype option on import or set low
_memory=False.

```
for chunk in pd.read_csv(file_path, chunksize=chunk_size):
```

C:\Users\User\AppData\Local\Temp\ipykernel_7652\1204666657.py:6: DtypeWarni
ng: Columns (5) have mixed types. Specify dtype option on import or set low
_memory=False.

```
for chunk in pd.read_csv(file_path, chunksize=chunk_size):
```

```
In [30]: labsstay.shape
```

```
Out[30]: (38372040, 20)
```

Common Labs

Average Observation Per Stay

getting the average number of times each itemid appears in an icustay

```
In [31]: labsstay_counts = labsstay['stay_id'].value_counts().reset_index()
labsstay_counts.columns = ['stay_id', 'labsstay_count']
labsstay_counts
```

```
Out[31]:
```

	stay_id	labsstay_count
0	36479755	63765
1	36671290	51252
2	31469106	48580
3	38199253	46222
4	34814635	46143
...
72093	36400520	1
72094	38046396	1
72095	32169536	1
72096	38886407	1
72097	39302677	1

72098 rows × 2 columns

```
In [32]: labsstay_counts.shape
```

```
Out[32]: (72098, 2)
```

```
In [33]: ICU_counts = icu_cohort['stay_id'].value_counts().reset_index()
ICU_counts.columns = ['stay_id', 'ICU_count']
ICU_counts
```

```
Out[33]:
```

	stay_id	ICU_count
0	31073147	37
1	34456715	37
2	32346798	37
3	35383104	37
4	34115393	37
...
72991	39824196	1
72992	33302469	1
72993	32622345	1
72994	33199830	1
72995	36195440	1

72996 rows × 2 columns

```
In [34]: ICU_counts.shape
```

```
Out[34]: (72996, 2)
```

```
In [35]: #keep stayid where labsstay count > icu count
merged_counts = ICU_counts.merge(labsstay_counts, on='stay_id', how='outer').ffill()
filtered_stayid = merged_counts[(merged_counts['labsstay_count']) > merged_counts['ICU_count']]
filtered_stayid
```

```
Out[35]:
```

	stay_id	ICU_count	labsstay_count
0	31073147	37	2701.0
1	34456715	37	370.0
2	32346798	37	111.0
3	35383104	37	2183.0
4	34115393	37	3663.0
...
72991	39824196	1	487.0
72992	33302469	1	621.0
72993	32622345	1	55.0
72994	33199830	1	24.0
72995	36195440	1	92.0

72042 rows × 3 columns

In [36]: `filtered_stayid.shape`

Out[36]: (72042, 3)

In [37]: `#Filter labsstays to only have these stay_ids`
`filtered_stayid_list = filtered_stayid['stay_id'].tolist()`

In [38]: `chunk_size = 100000`
`num_chunks = len(labsstay) // chunk_size + 1`

`chunks = []`

`# Process the DataFrame in chunks`
`for i in range(num_chunks):`
 `start_idx = i * chunk_size`
 `end_idx = (i + 1) * chunk_size`

`# Select the current chunk based on start and end indices`
`chunk = labsstay.iloc[start_idx:end_idx]`

`# Filter the chunk using 'filtered_stayid_list'`
`filtered_chunk = chunk[chunk['stay_id'].isin(filtered_stayid_list)]`

`# Append the processed chunk to the list`
`chunks.append(filtered_chunk)`

`# After processing all chunks, concatenate them into a final DataFrame`
`filtered_labsstay = pd.concat(chunks, ignore_index=True)`

In [39]: `filtered_labsstay.shape`

Out[39]: (38371885, 20)

In [40]: `# Group by 'itemid' and 'stay_id' and calculate the observation count for each`
`obs_per_stay = filtered_labsstay.groupby(['itemid', 'stay_id']).size().reset_index()`
`obs_per_stay.head()`

Out[40]:

	itemid	stay_id	count
0	50801	30004627	1
1	50801	30005362	1
2	50801	30007565	2
3	50801	30010930	1
4	50801	30011071	7

```
In [41]: min_count = obs_per_stay['count'].min()
max_count = obs_per_stay['count'].max()
avg_count = obs_per_stay['count'].mean()

print(f"Minimum Count: {min_count}")
print(f"Maximum Count: {max_count}")
print(f"Average Count: {avg_count}")
```

```
Minimum Count: 1
Maximum Count: 2340
Average Count: 11.43886010517153
```

```
In [42]: # Group by 'itemid' and calculate the average count for each 'itemid'
avg_obs_per_itemid = obs_per_stay.groupby('itemid')['count'].mean().reset_index()
avg_obs_per_itemid.head()
```

```
Out[42]:
```

	itemid	avg_obs
0	50801	3.885140
1	50802	20.290459
2	50803	3.004698
3	50804	20.291189
4	50805	2.327273

```
In [43]: min_count = avg_obs_per_itemid['avg_obs'].min()
max_count = avg_obs_per_itemid['avg_obs'].max()
avg_count = avg_obs_per_itemid['avg_obs'].mean()

print(f"Minimum Count: {min_count}")
print(f"Maximum Count: {max_count}")
print(f"Average Count: {avg_count}")
```

```
Minimum Count: 1.0
Maximum Count: 21.194129476701896
Average Count: 4.071320464236379
```

```
In [44]: # Filter the results based on the condition 'avg(count) > 10'
# we want the features to have more than 10 values entered for the average pat
avg_obs_per_stay = avg_obs_per_itemid[avg_obs_per_itemid['avg_obs'] > 10]
avg_obs_per_stay.head()
```

```
Out[44]:
```

	itemid	avg_obs
1	50802	20.290459
3	50804	20.291189
6	50808	15.252943
7	50809	12.308465
10	50813	12.994940

```
In [45]: avg_obs_per_stay.shape
```

```
Out[45]: (44, 2)
```

```
In [46]: # Merge 'labsstay' with 'd_labitems' on 'itemid'
#only keep the 44 itemid
chunk_size = 100000
num_chunks = len(filtered_labsstay) // chunk_size + 1

chunks = []

# Process the DataFrame in chunks
for i in range(num_chunks):
    start_idx = i * chunk_size
    end_idx = (i + 1) * chunk_size

    # Select the current chunk based on start and end indices
    chunk = filtered_labsstay.iloc[start_idx:end_idx]

    # Merge the current chunk with 'labitems' on 'itemid'
    merged_chunk = chunk.merge(labitems[['itemid', 'label']], on='itemid', how='left')
    # Merge the result with 'avg_obs_per_stay' on 'itemid'
    merged_chunk = merged_chunk.merge(avg_obs_per_stay[['avg_obs', 'itemid']], on='itemid', how='left')

    # Append the merged chunk to the list
    chunks.append(merged_chunk)

# After processing all chunks, concatenate them into a final DataFrame
df = pd.concat(chunks, ignore_index=True)
```

In [47]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32503253 entries, 0 to 32503252
Data columns (total 22 columns):
 #   Column                Dtype
---  -
 0   labevent_id           int64
 1   subject_id            int64
 2   hadm_id               float64
 3   specimen_id          int64
 4   itemid               int64
 5   order_provider_id    object
 6   charttime            datetime64[ns]
 7   storetime            object
 8   value                object
 9   valuenum             float64
10   valueuom             object
11   ref_range_lower      float64
12   ref_range_upper      float64
13   flag                 object
14   priority              object
15   comments              object
16   stay_id              int64
17   intime               datetime64[ns]
18   los                  float64
19   time_diff_days       float64
20   label                object
21   avg_obs              float64
dtypes: datetime64[ns](2), float64(7), int64(5), object(8)
memory usage: 5.3+ GB
```

In [48]: `# Group by 'Label' and 'avg_obs', and then calculate the count of distinct 'stay_id' for each group`
`commonlabs = df.groupby(['label', 'avg_obs']).agg(count=('stay_id', 'nunique'))`
`commonlabs.shape`

Out[48]: (44, 3)

In [49]: commonlabs

Out[49]:

	label	avg_obs	count
0	Amikacin	17.421053	76
1	Anion Gap	15.790314	71426
2	Base Excess	20.290459	44891
3	Bicarbonate	15.844876	71446
4	Calcium, Total	14.961132	68977
5	Calculated Total CO2	20.291189	44885
6	Chloride	16.497069	71475
7	Creatinine	15.829617	71480
8	Cyclosporin	14.471111	225
9	Free Calcium	15.252943	33213
10	Glucose	12.308465	24382
11	Glucose	15.615444	71381
12	H	19.829110	24741
13	Hematocrit	15.779609	71373
14	Hemoglobin	14.255044	71266
15	Heparin	18.980545	257
16	I	19.828981	24740
17	INR(PT)	10.830669	63798
18	L	19.829022	24740
19	Lactate	12.994940	46636
20	MCH	13.976786	71251
21	MCHC	13.980155	71252
22	MCV	13.977545	71254
23	Magnesium	15.514601	70338
24	Oxygen Saturation	13.402124	23734
25	PT	10.832409	63798
26	PTT	11.948269	63405
27	Phenobarbital	10.683706	313
28	Phosphate	15.047005	68992
29	Platelet Count	14.266822	71276
30	Potassium	16.777214	71481
31	Potassium, Whole Blood	11.162885	26092
32	RDW	13.971461	71236
33	RDW-SD	15.542799	30702
34	Rapamycin	11.573427	143
35	Red Blood Cells	13.977349	71256

	label	avg_obs	count
36	Sodium	16.693580	71477
37	Temperature	10.859185	19863
38	Urea Nitrogen	15.791723	71472
39	White Blood Cells	14.000000	71266
40	pCO2	20.290555	44883
41	pH	21.194129	46742
42	pO2	20.298839	44887
43	tacroFK	19.620936	1261

```
In [50]: commonlabs.to_csv('./data/preprocessing/commonlabs.csv.gz', index=False)
```

```
In [51]: del commonlabs
```

Labs Patients

extract the most common lab tests and the corresponding counts of how many patients have values for those labs

```
In [6]: commonlabs = pd.read_csv('./data/preprocessing/commonlabs.csv.gz', compression=
```

```

In [7]: file_path = './mimiciv/2.0/hosp/labevents.csv.gz'

chunk_size = 100000
chunks = []

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

    merged_chunk = pd.merge(chunk, labitems, on = 'itemid', how='inner')
    #merged_chunk = pd.merge(merged_chunk, commonlabs, on='label', how='inner')
    merged_chunk = pd.merge(merged_chunk, icu_cohort, on=['hadm_id', 'subject_id'])

    chunks.append(merged_chunk)

df = pd.concat(chunks, ignore_index=True)

```

C:\Users\User\AppData\Local\Temp\ipykernel_15052\2847547336.py:6: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

C:\Users\User\AppData\Local\Temp\ipykernel_15052\2847547336.py:6: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

C:\Users\User\AppData\Local\Temp\ipykernel_15052\2847547336.py:6: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

C:\Users\User\AppData\Local\Temp\ipykernel_15052\2847547336.py:6: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

C:\Users\User\AppData\Local\Temp\ipykernel_15052\2847547336.py:6: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

for chunk in pd.read_csv(file_path, chunksize=chunk_size):

```
In [8]: df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 102039133 entries, 0 to 102039132
Data columns (total 28 columns):
 #   Column                Dtype
---  -
 0   labevent_id           int64
 1   subject_id            int64
 2   hadm_id               float64
 3   specimen_id          int64
 4   itemid                int64
 5   order_provider_id     object
 6   charttime             object
 7   storetime             object
 8   value                 object
 9   valuenum              float64
10   valueuom              object
11   ref_range_lower       float64
12   ref_range_upper       float64
13   flag                  object
14   priority              object
15   comments              object
16   label                 object
17   fluid                 object
18   category              object
19   stay_id               int64
20   first_careunit        object
21   last_careunit         object
22   intime                object
23   outtime               object
24   los                   float64
25   hospital_expire_flag  int64
26   gender                object
27   anchor_age            int64
dtypes: float64(5), int64(7), object(16)
memory usage: 21.3+ GB
```

In [9]: `df.head()`

Out[9]:

	labevent_id	subject_id	hadm_id	specimen_id	itemid	order_provider_id	charttime	storetir
0	406	10000032	29079034.0	43001398	51237	NaN	2180-07-24 06:35:00	2180-C 08:10:
1	460	10000032	29079034.0	87246904	51237	NaN	2180-07-25 04:45:00	2180-C 07:07:
2	407	10000032	29079034.0	43001398	51274	NaN	2180-07-24 06:35:00	2180-C 08:10:
3	461	10000032	29079034.0	87246904	51274	NaN	2180-07-25 04:45:00	2180-C 07:07:
4	409	10000032	29079034.0	74547069	50861	NaN	2180-07-24 06:35:00	2180-C 08:11:

5 rows × 28 columns



In [10]: `df.shape`

Out[10]: (102039133, 28)

```
In [11]: #keep only the 44 common Labs
labels = commonlabs['label'].tolist()
chunk_size = 100000
num_chunks = len(df) // chunk_size + 1

chunks = []

# Process the DataFrame in chunks
for i in range(num_chunks):
    start_idx = i * chunk_size
    end_idx = (i + 1) * chunk_size

    chunk = df.iloc[start_idx:end_idx]

    filtered_chunk = chunk.merge(pd.DataFrame({'label': labels}), on='label',
    chunks.append(filtered_chunk)

filtered_df = pd.concat(chunks, ignore_index=True)
```

In [12]: `del df`

In [13]: `filtered_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 83646249 entries, 0 to 83646248
Data columns (total 28 columns):
#   Column                Dtype
---  -
0   labevent_id           int64
1   subject_id            int64
2   hadm_id               float64
3   specimen_id           int64
4   itemid                int64
5   order_provider_id     object
6   charttime             object
7   storetime             object
8   value                 object
9   valuenum              float64
10  valueuom              object
11  ref_range_lower        float64
12  ref_range_upper        float64
13  flag                   object
14  priority               object
15  comments               object
16  label                  object
17  fluid                  object
18  category               object
19  stay_id                int64
20  first_careunit         object
21  last_careunit          object
22  intime                 object
23  outtime                object
24  los                    float64
25  hospital_expire_flag   int64
26  gender                 object
27  anchor_age             int64
dtypes: float64(5), int64(7), object(16)
memory usage: 17.4+ GB
```

In [14]: `filtered_df.shape`

Out[14]: (83646249, 28)

In [15]: `import math`

In [16]: `filtered_df['charttime'] = pd.to_datetime(filtered_df['charttime'])`

In [17]: `filtered_df['intime'] = pd.to_datetime(filtered_df['intime'])`

In [18]: `# we want to extract measurements between admission and the end of the patient's stay`
`filtered_df['labresultoffset_min'] = ((filtered_df['charttime'] - filtered_df['intime'])`

```
In [19]: # we want to extract measurements between admission and the end of the patient's stay
filtered_df['labresultoffset_sec'] = (filtered_df['charttime'] - filtered_df['admissiontime'])
```

```
In [20]: filtered_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 83646249 entries, 0 to 83646248
Data columns (total 30 columns):
 #   Column                                Dtype
 ---  -
 0   labevent_id                          int64
 1   subject_id                          int64
 2   hadm_id                             float64
 3   specimen_id                         int64
 4   itemid                             int64
 5   order_provider_id                   object
 6   charttime                          datetime64[ns]
 7   storetime                          object
 8   value                              object
 9   valuenum                           float64
10   valueuom                           object
11   ref_range_lower                     float64
12   ref_range_upper                     float64
13   flag                               object
14   priority                           object
15   comments                           object
16   label                              object
17   fluid                              object
18   category                           object
19   stay_id                            int64
20   first_careunit                      object
21   last_careunit                       object
22   intime                             datetime64[ns]
23   outtime                            object
24   los                                float64
25   hospital_expire_flag                int64
26   gender                             object
27   anchor_age                         int64
28   labresultoffset_min                 int64
29   labresultoffset_sec                 float64
dtypes: datetime64[ns](2), float64(6), int64(8), object(14)
memory usage: 18.7+ GB
```

```
In [21]: # Keep only the rows where 'labresultoffset_sec' is between -1 and 'los' and no
chunk_size = 100000
num_chunks = len(filtered_df) // chunk_size + 1

chunks = []

for i in range(num_chunks):
    start_idx = i * chunk_size
    end_idx = (i + 1) * chunk_size

    chunk = filtered_df.iloc[start_idx:end_idx]
    filtered_chunk = chunk[(chunk['labresultoffset_sec'] >= -1) & (chunk['labr

    chunks.append(filtered_chunk)

lab = pd.concat(chunks, ignore_index=True)
```


In [22]: lab.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34200597 entries, 0 to 34200596
Data columns (total 30 columns):
#   Column                                Dtype
---  -
0   labevent_id                           int64
1   subject_id                           int64
2   hadm_id                               float64
3   specimen_id                           int64
4   itemid                               int64
5   order_provider_id                     object
6   charttime                             datetime64[ns]
7   storetime                             object
8   value                                 object
9   valuenum                              float64
10  valueuom                              object
11  ref_range_lower                        float64
12  ref_range_upper                        float64
13  flag                                   object
14  priority                               object
15  comments                              object
16  label                                 object
17  fluid                                 object
18  category                              object
19  stay_id                               int64
20  first_careunit                         object
21  last_careunit                         object
22  intime                                datetime64[ns]
23  outtime                               object
24  los                                    float64
25  hospital_expire_flag                  int64
26  gender                                object
27  anchor_age                            int64
28  labresultoffset_min                   int64
29  labresultoffset_sec                   float64
dtypes: datetime64[ns](2), float64(6), int64(8), object(14)
memory usage: 7.6+ GB
```

In [26]: lab.to_csv('./data/preprocessing/lab.csv.gz', index=False)

Chart Stay

extract the most common chartevents and the corresponding counts of how many patients have values for those chartevents

```

In [5]: #extracting the itemids for all the chartevents that occur within the time bound
chartevent_path = r"./mimiciv/2.0/icu/chartevents.csv.gz"
icucohort_path = r"./data/preprocessing/icu_cohort.csv.gz"

chunk_size = 100000

icucohort_df = pd.read_csv(icucohort_path)
chartevent_iter = pd.read_csv(chartevent_path, chunksize=chunk_size)

merged_chunks = []

# Iterate over the chunks and merge them with 'icucohort_df'
for chartevent_chunk in chartevent_iter:
    # Filter only numerical data from chartevent_chunk
    chartevent_chunk = chartevent_chunk[chartevent_chunk['valuenum'].notnull()]
    chartevent_chunk.drop(columns=['caregiver_id', 'warning', 'storetime'], inplace=True)

    # Merge chartevent_chunk with icucohort_df on 'stay_id'
    merged_chunk = pd.merge(chartevent_chunk, icucohort_df[['stay_id', 'intime', 'los']], on='stay_id', how='left')
    merged_chunks.append(merged_chunk)

    #convert intime and chartime to datetime
    merged_chunk['charttime'] = pd.to_datetime(merged_chunk['charttime'])
    merged_chunk['intime'] = pd.to_datetime(merged_chunk['intime'])

    #extract measurements between admission and the end of the patients' stay
    merged_chunk['time_difference_days'] = (merged_chunk['charttime'] - merged_chunk['intime']).dt.days

    # Filter the data where the time difference is between -1 and 'los'
    merged_chunk = merged_chunk[(merged_chunk['time_difference_days'] >= -1) & (merged_chunk['time_difference_days'] < merged_chunk['los'])]

    merged_chunk.drop(columns=['time_difference_days'], inplace=True)

# Concatenate the merged chunks into a single DataFrame
chartstay = pd.concat(merged_chunks)

# Save the merged DataFrame to a CSV file
chartstay.to_csv("./data/preprocessing/chartstay.csv.gz", index=False)

```

[s/stable/user_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
merged_chunk.drop(columns=['time_difference_days'], inplace=True)
```

C:\Users\User\AppData\Local\Temp\ipykernel_18476\163717513.py:32: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
merged_chunk.drop(columns=['time_difference_days'], inplace=True)
```

C:\Users\User\AppData\Local\Temp\ipykernel_18476\163717513.py:32: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

In [6]: `chartstay.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 295056465 entries, 0 to 21355
```

```
Data columns (total 11 columns):
```

#	Column	Dtype
0	subject_id	int64
1	hadm_id	int64
2	stay_id	int64
3	charttime	datetime64[ns]
4	itemid	int64
5	value	object
6	valuenum	float64
7	valueuom	object
8	intime	datetime64[ns]
9	los	float64
10	time_difference_days	float64

```
dtypes: datetime64[ns](2), float64(3), int64(4), object(2)
```

```
memory usage: 26.4+ GB
```

In [7]: `chartstay.shape`

Out[7]: (295056465, 11)

```

In [11]: #getting the average number of times each itemid appears in an icustay
columns_needed = ['itemid', 'stay_id']

chunksize = 100000
data_chunks = pd.read_csv("../data/preprocessing/chartstay.csv.gz", usecols=columns_needed, chunksize=chunksize)

obs_per_stay = pd.DataFrame()

for chunk in data_chunks:
    # calculate counts per stay_id for each itemid
    obs_per_stay_chunk = chunk.groupby(['itemid', 'stay_id']).size().reset_index(name='count')

    # Concatenate the results of each chunk to the main DataFrame
    obs_per_stay = pd.concat([obs_per_stay, obs_per_stay_chunk])

# Calculate the average number of observations per stay for each itemid
avg_obs_per_stay = obs_per_stay.groupby('itemid')['count'].mean().reset_index(name='avg_obs')

avg_obs_per_stay.describe()

```

Out[11]:

	itemid	avg_obs
count	939.000000	939.000000
mean	226524.889244	48.714249
std	2366.602913	104.407604
min	220045.000000	1.000000
25%	224900.500000	5.082748
50%	226998.000000	14.450213
75%	228233.500000	43.906732
max	229882.000000	976.492958

In [12]: avg_obs_per_stay

Out[12]:

	itemid	avg_obs
0	220045	208.922726
1	220046	20.095912
2	220047	20.104746
3	220050	172.164118
4	220051	172.098523
...
934	229865	2.037221
935	229872	9.250000
936	229880	39.590476
937	229881	65.483871
938	229882	28.722222

939 rows × 2 columns

In [13]: avg_obs_per_stay.shape

Out[13]: (939, 2)

In [14]: *# Filter the results to keep only those with an average count greater than 49*
 avg_obs_per_stay = avg_obs_per_stay[avg_obs_per_stay['avg_obs'] > 49]

In [15]: avg_obs_per_stay.shape

Out[15]: (199, 2)

Common Charts

extract the most common chartevents and the corresponding counts of how many patients have values for those chartevents

In [7]: d_items = pd.read_csv('./mimiciv/2.0/icu/d_items.csv.gz',compression='gzip')

In [8]: `d_items.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4014 entries, 0 to 4013
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   itemid                4014 non-null   int64
1   label                 4014 non-null   object
2   abbreviation          4014 non-null   object
3   linksto               4014 non-null   object
4   category              4014 non-null   object
5   unitname             1592 non-null   object
6   param_type            4014 non-null   object
7   lownormalvalue        19 non-null     float64
8   highnormalvalue       22 non-null     float64
dtypes: float64(2), int64(1), object(6)
memory usage: 282.4+ KB
```

In [9]: `d_items.head()`

Out[9]:

	itemid	label	abbreviation	linksto	category	unitname	param_type	lownormalval
0	220001	Problem List	Problem List	chartevents	General	NaN	Text	N
1	220003	ICU Admission date	ICU Admission date	datetimeevents	ADT	NaN	Date and time	N
2	220045	Heart Rate	HR	chartevents	Routine Vital Signs	bpm	Numeric	N
3	220046	Heart rate Alarm - High	HR Alarm - High	chartevents	Alarms	bpm	Numeric	N
4	220047	Heart Rate Alarm - Low	HR Alarm - Low	chartevents	Alarms	bpm	Numeric	N

```
In [16]: columns_needed = ['itemid', 'stay_id']

chunksize = 100000
data_chunks = pd.read_csv("./data/preprocessing/chartstay.csv.gz", usecols=columns_needed,
                           chunksize=chunksize)

result = pd.DataFrame()

for chunk in data_chunks:
    merged_data = chunk.merge(d_items, on='itemid', how='inner')
    merged_data = merged_data.merge(avg_obs_per_stay, on='itemid', how='inner')

    result = pd.concat([result, merged_data])
```

```
In [17]: result.to_csv("./data/preprocessing/merge_ditems_chart_avgobs.csv.gz", index=False)
```

```
In [7]: result = pd.read_csv('./data/preprocessing/merge_ditems_chart_avgobs.csv.gz')
```

```
In [8]: grouped_result = result.groupby(['label', 'avg_obs'])['stay_id'].nunique().reset_index()
```

```
In [9]: # Calculate the threshold for the minimum number of stay_ids required
threshold = len(icu_cohort['stay_id'].unique()) * 0.25
threshold
```

```
Out[9]: 18249.0
```

```
In [10]: # Filter the results to keep only rows where count is greater than the threshold
filtered_result = grouped_result[grouped_result['count'] > threshold]
```

```
In [11]: # Sort the filtered results by count in descending order
ld_commonchart = filtered_result.sort_values(by='count', ascending=False)
```

In [12]: `ld_commonchart`

Out[12]:

		label	avg_obs	count
75		Heart Rate	208.922726	72978
164		Respiratory Rate	207.225807	72907
108		O2 saturation pulseoxymetry	204.530713	72902
62		GCS - Eye Opening	49.738155	72572
64		GCS - Verbal Response	49.641869	72561
63		GCS - Motor Response	49.530978	72551
106		Non Invasive Blood Pressure mean	139.544098	72067
107		Non Invasive Blood Pressure systolic	139.534886	72061
105		Non Invasive Blood Pressure diastolic	139.505182	72058
186		Temperature Fahrenheit	50.333491	71800
92		Inspired O2 Fraction	54.795313	36455
7		Activity / Mobility (JH-HLM)	85.090176	31219
16		Arterial Blood Pressure mean	171.522424	27785
15		Arterial Blood Pressure diastolic	172.098523	27574
17		Arterial Blood Pressure systolic	172.164118	27566
109		Orientation	166.041201	23261

In [13]: `ld_commonchart.shape`

Out[13]: (16, 3)

In [14]: `ld_commonchart.to_csv("./data/preprocessing/ld_commonchart.csv.gz", index=False)`

Chartevents Processed

Keep features from the most common chart features

In [8]: `ld_commonchart= pd.read_csv('./data/preprocessing/ld_commonchart.csv.gz')`


```

In [11]: def calculate_chartoffset(charttime, intime):
          return (charttime - intime).total_seconds() // 60

columns_needed = ['stay_id', 'charttime', 'itemid', 'valuenum']

chunksize = 10000
data_chunks = pd.read_csv("./mimiciv/2.0/icu/chartevents.csv.gz", usecols=columns_needed)

ld = pd.DataFrame()

for chunk in data_chunks:

    # Join chunk with d_items on itemid
    merged_data = chunk.merge(d_items, on='itemid', how='inner')

    # Join merged_data with ld_commonchart on label
    merged_data = merged_data.merge(ld_commonchart, on='label', how='inner')

    # Join merged_data with ld_labels on stay_id
    merged_data = merged_data.merge(icu_cohort[['stay_id', 'intime', 'los']], on='stay_id', how='inner')

    # Convert 'charttime' & 'intime' column to datetime type
    merged_data['charttime'] = pd.to_datetime(merged_data['charttime'])
    merged_data['intime'] = pd.to_datetime(merged_data['intime'])

    # Calculate chartoffset using the calculate_chartoffset function
    merged_data['chartoffset'] = merged_data.apply(lambda row: calculate_chartoffset(row['charttime'], row['intime']), axis=1)

    # Filter data based on the time interval and non-null valuenum
    filtered_data = merged_data[(merged_data['chartoffset'] >= -60) & (merged_data['valuenum'].notnull())]

    # Concatenate the processed chunk to the ld DataFrame
    ld = pd.concat([ld, filtered_data])

In [13]: ld.to_csv("./data/preprocessing/ld.csv.gz", index=False)
          #ld = pd.read_csv('./data/preprocessing/ld.csv.gz')

```

In [15]: `ld.head()`

Out[15]:

	stay_id	charttime	itemid	valuenum	label	abbreviation	linksto	category	unitnam
0	39553978	2180-07-23 21:01:00	220179	82.0	Non Invasive Blood Pressure systolic	NBPs	chartevents	Routine Vital Signs	mmH
1	39553978	2180-07-23 22:00:00	220179	85.0	Non Invasive Blood Pressure systolic	NBPs	chartevents	Routine Vital Signs	mmH
2	39553978	2180-07-23 19:00:00	220179	93.0	Non Invasive Blood Pressure systolic	NBPs	chartevents	Routine Vital Signs	mmH
3	39553978	2180-07-23 20:00:00	220179	90.0	Non Invasive Blood Pressure systolic	NBPs	chartevents	Routine Vital Signs	mmH
4	39553978	2180-07-23 14:11:00	220179	84.0	Non Invasive Blood Pressure systolic	NBPs	chartevents	Routine Vital Signs	mmH

Combine Data

```
In [5]: def process_mimiciv(input_file, output_file, columns_to_include):
# Read the entire CSV file
data = pd.read_csv(input_file)

# Calculate the average valuenum for each distinct stay_id and itemid combination
avg_valuenum = data.groupby(['stay_id', 'label'])['valuenum'].mean().reset_index()

# Pivot the data to create distinct itemid columns
pivoted_data = avg_valuenum.pivot(index='stay_id', columns='label',
                                  values='valuenum').reset_index()

# Replace NaN values with 0 in every column
pivoted_data = pivoted_data.fillna(0)

# Select the specified columns to include in the final processed data
final_data = pivoted_data.merge(data[columns_to_include].drop_duplicates(),
                                how='inner')

# Save the processed data to a CSV file
final_data.to_csv(output_file, index=False)
```

```
In [6]: columns_to_include = ['stay_id', 'los']
process_mimiciv("./data/preprocessing/ld.csv.gz", "./data/preprocessing/processed_chart.csv.gz",
               columns_to_include)
```

```
In [7]: processed_chart= pd.read_csv('./data/preprocessing/processed_chart.csv.gz')
```

```
In [8]: processed_chart.head()
```

```
Out[8]:
```

	stay_id	Activity / Mobility (JH-HLM)	Arterial Blood Pressure diastolic	Arterial Blood Pressure mean	Arterial Blood Pressure systolic	GCS - Eye Opening	GCS - Motor Response	GCS - Verbal Response	Heart F
0	30000153	0.0	66.935484	90.300000	137.387097	3.541667	5.875000	3.458333	104.804
1	30000213	3.5	0.000000	0.000000	0.000000	3.818182	5.818182	3.636364	82.736
2	30000484	0.0	0.000000	0.000000	0.000000	3.866667	4.933333	3.066667	91.881
3	30000646	0.0	0.000000	0.000000	0.000000	4.000000	6.000000	5.000000	94.046
4	30001148	0.0	58.833333	73.541667	108.291667	3.285714	5.285714	3.857143	74.592

```
In [9]: processed_chart.shape
```

```
Out[9]: (72980, 18)
```

```
In [10]: columns_to_include = ['stay_id', 'los', 'gender', 'anchor_age']
process_mimiciv("./data/preprocessing/lab.csv.gz", "./data/preprocessing/processed_lab.csv.gz",
               columns_to_include)
```

C:\Users\User\AppData\Local\Temp\ipykernel_1424\2272975284.py:3: DtypeWarning: Columns (5) have mixed types. Specify dtype option on import or set low_memory=False.

```
data = pd.read_csv(input_file)
```

```
In [11]: processed_lab= pd.read_csv('./data/preprocessing/processed_lab.csv.gz')
```

```
In [12]: processed_lab.head()
```

```
Out[12]:
```

	stay_id	Amikacin	Anion Gap	Base Excess	Bicarbonate	Calcium, Total	Calculated Total CO2	Chloride	Cre
0	30000153	0.0	12.000000	-3.333333	21.000000	7.700000	22.666667	115.000000	1.0
1	30000213	0.0	15.666667	0.500000	23.333333	8.333333	27.000000	100.666667	3.0
2	30000484	0.0	10.000000	1.000000	27.666667	8.133333	33.000000	105.000000	1.2
3	30000646	0.0	12.000000	-1.000000	22.125000	7.625000	21.000000	109.500000	0.7
4	30001148	0.0	10.500000	2.444444	27.666667	0.000000	28.111111	106.666667	0.7

5 rows × 47 columns

```
In [13]: processed_lab.shape
```

```
Out[13]: (72043, 47)
```


```
In [14]: processed_data = processed_lab.merge(processed_chart, on=['stay_id', 'los'], how='outer')
```

```
In [15]: processed_data.head()
```

```
Out[15]:
```

	stay_id	Amikacin	Anion Gap	Base Excess	Bicarbonate	Calcium, Total	Calculated Total CO2	Chloride	Creatinine
0	30000153	0.0	12.000000	-3.333333	21.000000	7.700000	22.666667	115.000000	1.000000
1	30000213	0.0	15.666667	0.500000	23.333333	8.333333	27.000000	100.666667	3.000000
2	30000484	0.0	10.000000	1.000000	27.666667	8.133333	33.000000	105.000000	1.500000
3	30000646	0.0	12.000000	-1.000000	22.125000	7.625000	21.000000	109.500000	0.750000
4	30001148	0.0	10.500000	2.444444	27.666667	0.000000	28.111111	106.666667	0.750000

5 rows × 63 columns



```
In [16]: processed_data.shape
```

```
Out[16]: (72030, 63)
```

```
In [17]: processed_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 72030 entries, 0 to 72029
```

```
Data columns (total 63 columns):
```

#	Column	Non-Null Count	Dtype
0	stay_id	72030 non-null	int64
1	Amikacin	72030 non-null	float64
2	Anion Gap	72030 non-null	float64
3	Base Excess	72030 non-null	float64
4	Bicarbonate	72030 non-null	float64
5	Calcium, Total	72030 non-null	float64
6	Calculated Total CO2	72030 non-null	float64
7	Chloride	72030 non-null	float64
8	Creatinine	72030 non-null	float64
9	Cyclosporin	72030 non-null	float64
10	Free Calcium	72030 non-null	float64
11	Glucose	72030 non-null	float64
12	H	72030 non-null	float64
13	Hematocrit	72030 non-null	float64
14	Hemoglobin	72030 non-null	float64
15	Heparin	72030 non-null	float64
16	I	72030 non-null	float64
17	INR(PT)	72030 non-null	float64
18	L	72030 non-null	float64
19	Lactate	72030 non-null	float64
20	MCH	72030 non-null	float64
21	MCHC	72030 non-null	float64
22	MCV	72030 non-null	float64
23	Magnesium	72030 non-null	float64
24	Oxygen Saturation	72030 non-null	float64
25	PT	72030 non-null	float64
26	PTT	72030 non-null	float64
27	Phenobarbital	72030 non-null	float64
28	Phosphate	72030 non-null	float64
29	Platelet Count	72030 non-null	float64
30	Potassium	72030 non-null	float64
31	Potassium, Whole Blood	72030 non-null	float64
32	RDW	72030 non-null	float64
33	RDW-SD	72030 non-null	float64
34	Rapamycin	72030 non-null	float64
35	Red Blood Cells	72030 non-null	float64
36	Sodium	72030 non-null	float64
37	Temperature	72030 non-null	float64
38	Urea Nitrogen	72030 non-null	float64
39	White Blood Cells	72030 non-null	float64
40	pCO2	72030 non-null	float64
41	pH	72030 non-null	float64
42	pO2	72030 non-null	float64
43	tacroFK	72030 non-null	float64
44	los	72030 non-null	float64
45	gender	72030 non-null	object
46	anchor_age	72030 non-null	int64
47	Activity / Mobility (JH-HLM)	72030 non-null	float64
48	Arterial Blood Pressure diastolic	72030 non-null	float64
49	Arterial Blood Pressure mean	72030 non-null	float64
50	Arterial Blood Pressure systolic	72030 non-null	float64
51	GCS - Eye Opening	72030 non-null	float64

```

52 GCS - Motor Response          72030 non-null float64
53 GCS - Verbal Response         72030 non-null float64
54 Heart Rate                    72030 non-null float64
55 Inspired O2 Fraction          72030 non-null float64
56 Non Invasive Blood Pressure diastolic 72030 non-null float64
57 Non Invasive Blood Pressure mean  72030 non-null float64
58 Non Invasive Blood Pressure systolic 72030 non-null float64
59 O2 saturation pulseoxymetry     72030 non-null float64
60 Orientation                   72030 non-null float64
61 Respiratory Rate              72030 non-null float64
62 Temperature Fahrenheit        72030 non-null float64
dtypes: float64(60), int64(2), object(1)
memory usage: 35.2+ MB

```

```
In [19]: processed_data.to_csv('./data/preprocessing/processed_data.csv.gz', index=False)
```


```
In [20]: # remove some binary and less useful variables from the original set
los=processed_data.drop(columns=['stay_id', 'Orientation', 'Temperature', 'Base (mEq/L)',
                                'Amikacin', 'Cyclosporin'])
```

```
In [21]: los.tail()
```

Out[21]:

	Anion Gap	Bicarbonate	Calcium, Total	Calculated Total CO2	Chloride	Creatinine	Free Calcium	Glucose
72025	13.500000	23.500000	8.80	22.666667	108.500000	1.65	1.1200	184.000000
72026	16.000000	24.000000	8.60	0.000000	103.000000	2.20	1.1100	116.000000
72027	11.500000	26.000000	8.80	27.166667	105.666667	0.70	1.1275	114.444444
72028	11.333333	24.333333	8.50	0.000000	103.666667	0.90	0.0000	106.000000
72029	15.400000	23.400000	8.88	0.000000	104.400000	1.00	0.0000	112.000000

5 rows × 57 columns



In [22]: `los.info()`


```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 72030 entries, 0 to 72029
```

```
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Anion Gap	72030 non-null	float64
1	Bicarbonate	72030 non-null	float64
2	Calcium, Total	72030 non-null	float64
3	Calculated Total CO2	72030 non-null	float64
4	Chloride	72030 non-null	float64
5	Creatinine	72030 non-null	float64
6	Free Calcium	72030 non-null	float64
7	Glucose	72030 non-null	float64
8	H	72030 non-null	float64
9	Hematocrit	72030 non-null	float64
10	Hemoglobin	72030 non-null	float64
11	Heparin	72030 non-null	float64
12	I	72030 non-null	float64
13	INR(PT)	72030 non-null	float64
14	L	72030 non-null	float64
15	Lactate	72030 non-null	float64
16	MCH	72030 non-null	float64
17	MCHC	72030 non-null	float64
18	MCV	72030 non-null	float64
19	Magnesium	72030 non-null	float64
20	Oxygen Saturation	72030 non-null	float64
21	PT	72030 non-null	float64
22	PTT	72030 non-null	float64
23	Phenobarbital	72030 non-null	float64
24	Phosphate	72030 non-null	float64
25	Platelet Count	72030 non-null	float64
26	Potassium	72030 non-null	float64
27	Potassium, Whole Blood	72030 non-null	float64
28	RDW	72030 non-null	float64
29	RDW-SD	72030 non-null	float64
30	Rapamycin	72030 non-null	float64
31	Red Blood Cells	72030 non-null	float64
32	Sodium	72030 non-null	float64
33	Urea Nitrogen	72030 non-null	float64
34	White Blood Cells	72030 non-null	float64
35	pCO2	72030 non-null	float64
36	pH	72030 non-null	float64
37	pO2	72030 non-null	float64
38	tacroFK	72030 non-null	float64
39	los	72030 non-null	float64
40	gender	72030 non-null	object
41	anchor_age	72030 non-null	int64
42	Activity / Mobility (JH-HLM)	72030 non-null	float64
43	Arterial Blood Pressure diastolic	72030 non-null	float64
44	Arterial Blood Pressure mean	72030 non-null	float64
45	Arterial Blood Pressure systolic	72030 non-null	float64
46	GCS - Eye Opening	72030 non-null	float64
47	GCS - Motor Response	72030 non-null	float64
48	GCS - Verbal Response	72030 non-null	float64
49	Heart Rate	72030 non-null	float64
50	Inspired O2 Fraction	72030 non-null	float64
51	Non Invasive Blood Pressure diastolic	72030 non-null	float64

```
52 Non Invasive Blood Pressure mean      72030 non-null float64
53 Non Invasive Blood Pressure systolic  72030 non-null float64
54 O2 saturation pulseoxymetry          72030 non-null float64
55 Respiratory Rate                      72030 non-null float64
56 Temperature Fahrenheit                72030 non-null float64
dtypes: float64(55), int64(1), object(1)
memory usage: 31.9+ MB
```

```
In [23]: los['gender'].replace({'M': 0, 'F':1}, inplace=True)
```

In [24]: `los.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 72030 entries, 0 to 72029
```

```
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
0	Anion Gap	72030 non-null	float64
1	Bicarbonate	72030 non-null	float64
2	Calcium, Total	72030 non-null	float64
3	Calculated Total CO2	72030 non-null	float64
4	Chloride	72030 non-null	float64
5	Creatinine	72030 non-null	float64
6	Free Calcium	72030 non-null	float64
7	Glucose	72030 non-null	float64
8	H	72030 non-null	float64
9	Hematocrit	72030 non-null	float64
10	Hemoglobin	72030 non-null	float64
11	Heparin	72030 non-null	float64
12	I	72030 non-null	float64
13	INR(PT)	72030 non-null	float64
14	L	72030 non-null	float64
15	Lactate	72030 non-null	float64
16	MCH	72030 non-null	float64
17	MCHC	72030 non-null	float64
18	MCV	72030 non-null	float64
19	Magnesium	72030 non-null	float64
20	Oxygen Saturation	72030 non-null	float64
21	PT	72030 non-null	float64
22	PTT	72030 non-null	float64
23	Phenobarbital	72030 non-null	float64
24	Phosphate	72030 non-null	float64
25	Platelet Count	72030 non-null	float64
26	Potassium	72030 non-null	float64
27	Potassium, Whole Blood	72030 non-null	float64
28	RDW	72030 non-null	float64
29	RDW-SD	72030 non-null	float64
30	Rapamycin	72030 non-null	float64
31	Red Blood Cells	72030 non-null	float64
32	Sodium	72030 non-null	float64
33	Urea Nitrogen	72030 non-null	float64
34	White Blood Cells	72030 non-null	float64
35	pCO2	72030 non-null	float64
36	pH	72030 non-null	float64
37	pO2	72030 non-null	float64
38	tacroFK	72030 non-null	float64
39	los	72030 non-null	float64
40	gender	72030 non-null	int64
41	anchor_age	72030 non-null	int64
42	Activity / Mobility (JH-HLM)	72030 non-null	float64
43	Arterial Blood Pressure diastolic	72030 non-null	float64
44	Arterial Blood Pressure mean	72030 non-null	float64
45	Arterial Blood Pressure systolic	72030 non-null	float64
46	GCS - Eye Opening	72030 non-null	float64
47	GCS - Motor Response	72030 non-null	float64
48	GCS - Verbal Response	72030 non-null	float64
49	Heart Rate	72030 non-null	float64
50	Inspired O2 Fraction	72030 non-null	float64
51	Non Invasive Blood Pressure diastolic	72030 non-null	float64

```
52 Non Invasive Blood Pressure mean      72030 non-null float64
53 Non Invasive Blood Pressure systolic  72030 non-null float64
54 O2 saturation pulseoxymetry           72030 non-null float64
55 Respiratory Rate                       72030 non-null float64
56 Temperature Fahrenheit                 72030 non-null float64
dtypes: float64(55), int64(2)
memory usage: 31.9 MB
```

```
In [28]: los.to_csv('./data/dataset/los.csv.gz', index=False)
```

Length-of-Stay Prediction Model

```
In [3]: df = pd.read_csv('./data/dataset/los.csv.gz')
```

In [4]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 72030 entries, 0 to 72029
```

```
Data columns (total 57 columns):
```

#	Column	Non-Null Count	Dtype
0	Anion Gap	72030 non-null	float64
1	Bicarbonate	72030 non-null	float64
2	Calcium, Total	72030 non-null	float64
3	Calculated Total CO2	72030 non-null	float64
4	Chloride	72030 non-null	float64
5	Creatinine	72030 non-null	float64
6	Free Calcium	72030 non-null	float64
7	Glucose	72030 non-null	float64
8	H	72030 non-null	float64
9	Hematocrit	72030 non-null	float64
10	Hemoglobin	72030 non-null	float64
11	Heparin	72030 non-null	float64
12	I	72030 non-null	float64
13	INR(PT)	72030 non-null	float64
14	L	72030 non-null	float64
15	Lactate	72030 non-null	float64
16	MCH	72030 non-null	float64
17	MCHC	72030 non-null	float64
18	MCV	72030 non-null	float64
19	Magnesium	72030 non-null	float64
20	Oxygen Saturation	72030 non-null	float64
21	PT	72030 non-null	float64
22	PTT	72030 non-null	float64
23	Phenobarbital	72030 non-null	float64
24	Phosphate	72030 non-null	float64
25	Platelet Count	72030 non-null	float64
26	Potassium	72030 non-null	float64
27	Potassium, Whole Blood	72030 non-null	float64
28	RDW	72030 non-null	float64
29	RDW-SD	72030 non-null	float64
30	Rapamycin	72030 non-null	float64
31	Red Blood Cells	72030 non-null	float64
32	Sodium	72030 non-null	float64
33	Urea Nitrogen	72030 non-null	float64
34	White Blood Cells	72030 non-null	float64
35	pCO2	72030 non-null	float64
36	pH	72030 non-null	float64
37	pO2	72030 non-null	float64
38	tacroFK	72030 non-null	float64
39	los	72030 non-null	float64
40	gender	72030 non-null	int64
41	anchor_age	72030 non-null	int64
42	Activity / Mobility (JH-HLM)	72030 non-null	float64
43	Arterial Blood Pressure diastolic	72030 non-null	float64
44	Arterial Blood Pressure mean	72030 non-null	float64
45	Arterial Blood Pressure systolic	72030 non-null	float64
46	GCS - Eye Opening	72030 non-null	float64
47	GCS - Motor Response	72030 non-null	float64
48	GCS - Verbal Response	72030 non-null	float64
49	Heart Rate	72030 non-null	float64
50	Inspired O2 Fraction	72030 non-null	float64
51	Non Invasive Blood Pressure diastolic	72030 non-null	float64

```

52 Non Invasive Blood Pressure mean      72030 non-null float64
53 Non Invasive Blood Pressure systolic  72030 non-null float64
54 O2 saturation pulseoxymetry           72030 non-null float64
55 Respiratory Rate                       72030 non-null float64
56 Temperature Fahrenheit                 72030 non-null float64
dtypes: float64(55), int64(2)
memory usage: 31.3 MB

```

```

In [5]: # Target Variable (Length-of-Stay)
        LOS = df['los'].values
        # Prediction Features
        features = df.drop(columns=['los'])

```

```

In [6]: actual_mean_los = np.mean(LOS)
        actual_median_los = np.median(LOS)
        print(f"actual mean LOS = {actual_mean_los}\nactual median los = {actual_median_los}")

actual mean LOS = 3.493232738352332
actual median los = 1.9532812499999999

```

```

In [7]: # Split into train 80% and test 20%
        X_train, X_test, y_train, y_test = train_test_split(features,
                                                             LOS,
                                                             test_size = .20,
                                                             random_state = 0)

        # Show the results of the split
        print("Training set has {} samples.".format(X_train.shape[0]))
        print("Testing set has {} samples.".format(X_test.shape[0]))

Training set has 57624 samples.
Testing set has 14406 samples.

```



```
In [19]: from sklearn.model_selection import cross_val_score
# Regression models for comparison
models = [GradientBoostingRegressor(random_state=0),
          KNeighborsRegressor(),
          RandomForestRegressor(random_state=0),
          XGBRegressor(random_state=0)]

mae = {}
mse = {}
rmse = {}
r2 = {}

for model in models:
    # Instantiate and fit Regressor Model
    reg_model = model
    reg_model.fit(X_train, y_train)

    # Make predictions with model
    y_test_preds = reg_model.predict(X_test)

    # Grab model name and store results associated with model
    name = str(model).split("(")[0]

    # Evaluation metrics
    mae[name] = mean_absolute_error(y_test, y_test_preds)
    mse[name] = mean_squared_error(y_test, y_test_preds)
    rmse[name] = mean_squared_error(y_test, y_test_preds, squared=False)
    r2[name] = r2_score(y_test, y_test_preds)

    print('{} done.'.format(name))
```

GradientBoostingRegressor done.

KNeighborsRegressor done.


RandomForestRegressor done.

XGBRegressor done.

Model Evaluation

```
In [20]: # R2 score results
fig, ax = plt.subplots()
ind = range(len(r2))
ax.barh(ind, list(r2.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(r2.keys())
ax.set_xlabel('R-squared score')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
#fig.savefig('images/compare_models.png', bbox_inches = 'tight')
```

Out[20]: Text(0.5, 1.0, 'Comparison of Regression Models')




```
In [21]: #R2 Values
for key, value in r2.items():
    print(f"Key: {key}, Value: {value}")
```

Key: GradientBoostingRegressor, Value: 0.520799935008655
Key: KNeighborsRegressor, Value: 0.2196412208098515
Key: RandomForestRegressor, Value: 0.5284204642419412
Key: XGBRegressor, Value: 0.604343214448636

```
In [22]: # mse score results
fig, ax = plt.subplots()
ind = range(len(mse))
ax.barh(ind, list(mse.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(mse.keys())
ax.set_xlabel('Mean Squared Error (MSE)')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
```

Out[22]: Text(0.5, 1.0, 'Comparison of Regression Models')

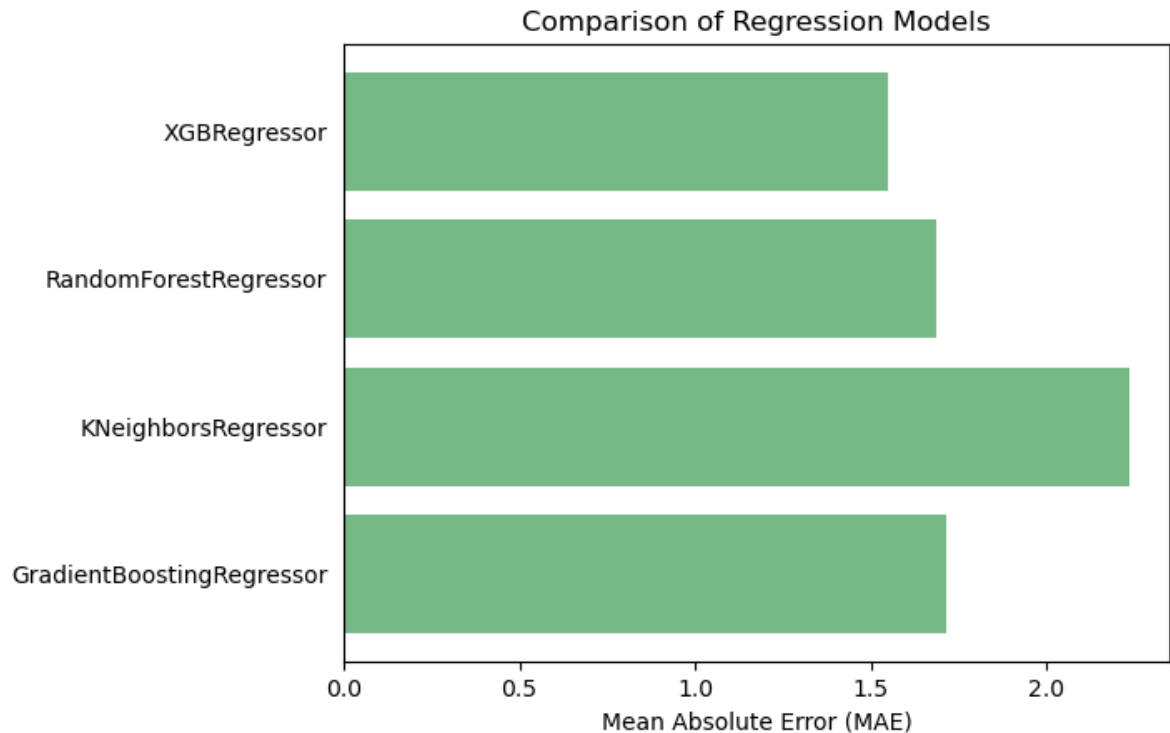


```
In [23]: #MSE Values
for key, value in mse.items():
    print(f"Key: {key}, Value: {value}")
```

Key: GradientBoostingRegressor, Value: 10.493854867362673
Key: KNeighborsRegressor, Value: 17.08883694212695
Key: RandomForestRegressor, Value: 10.326975241025298
Key: XGBRegressor, Value: 8.664366280789716

```
In [24]: # mae score results
fig, ax = plt.subplots()
ind = range(len(mae))
ax.barh(ind, list(mae.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(mae.keys())
ax.set_xlabel('Mean Absolute Error (MAE)')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
```

Out[24]: Text(0.5, 1.0, 'Comparison of Regression Models')

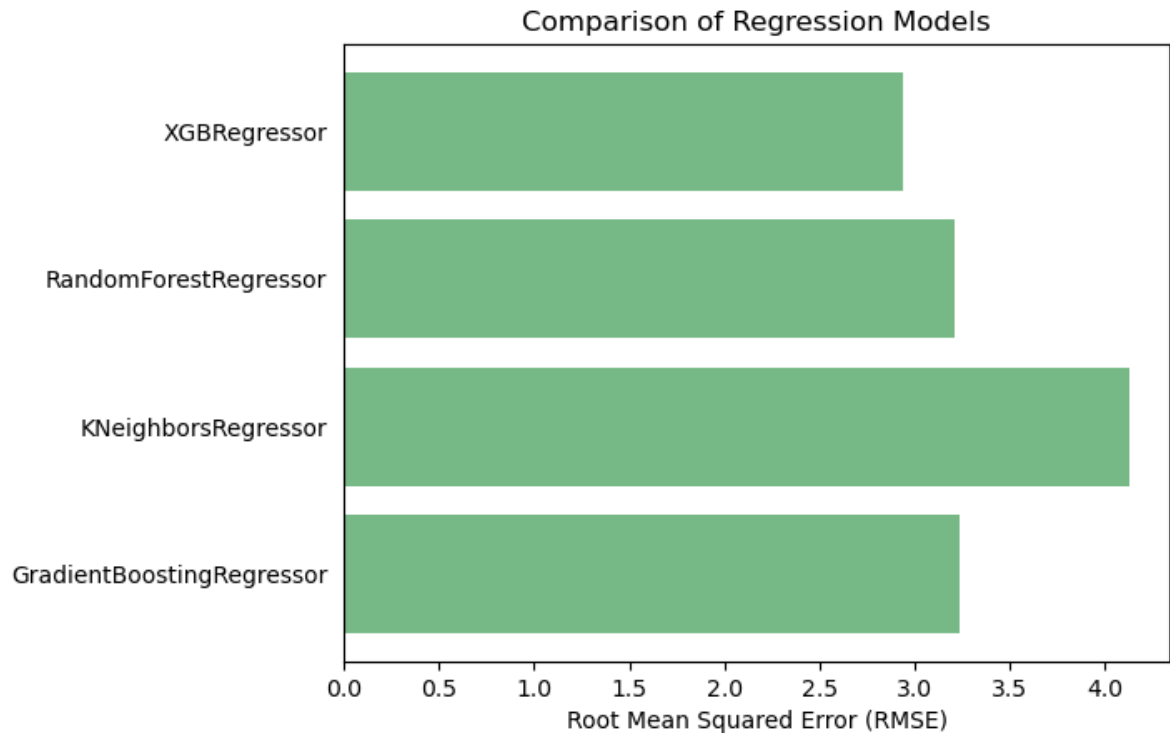


```
In [25]: #MAE Values
for key, value in mae.items():
    print(f"Key: {key}, Value: {value}")
```

Key: GradientBoostingRegressor, Value: 1.7147232102121142
Key: KNeighborsRegressor, Value: 2.238807677870846
Key: RandomForestRegressor, Value: 1.6877608309744911
Key: XGBRegressor, Value: 1.5476066202460712

```
In [26]: # rmse score results
fig, ax = plt.subplots()
ind = range(len(rmse))
ax.barh(ind, list(rmse.values()), align='center',
        color = '#55a868', alpha=0.8)
ax.set_yticks(ind)
ax.set_yticklabels(rmse.keys())
ax.set_xlabel('Root Mean Squared Error (RMSE)')
ax.tick_params(left=False, top=False, right=False)
ax.set_title('Comparison of Regression Models')
```

Out[26]: Text(0.5, 1.0, 'Comparison of Regression Models')



```
In [27]: #RMSE value
for key, value in rmse.items():
    print(f"Key: {key}, Value: {value}")
```

Key: GradientBoostingRegressor, Value: 3.2394219958756025
Key: KNeighborsRegressor, Value: 4.133864649710601
Key: RandomForestRegressor, Value: 3.2135611463025406
Key: XGBRegressor, Value: 2.9435295617319213

```
In [15]: # XGBRegressor will be used as the LOS prediction model
reg_model = XGBRegressor(random_state=0)
reg_model.fit(X_train, y_train)
y_test_preds = reg_model.predict(X_test)
r2_not_refined = r2_score(y_test, y_test_preds)
mse_not_refined = mean_squared_error(y_test, y_test_preds)
mae_not_refined = mean_absolute_error(y_test, y_test_preds)
rmse_not_refined = mean_squared_error(y_test, y_test_preds, squared=False)
print(f"R2 score is: {r2_not_refined}")
print(f"MSE score is: {mse_not_refined}")
print(f"MAE score is: {mae_not_refined}")
print(f"RMSE score is: {rmse_not_refined}")
```

R2 score is: 0.604343214448636

MSE score is: 8.664366280789716

MAE score is: 1.5476066202460712

RMSE score is: 2.9435295617319213

Model Refinement

```

In [10]: # Split into train 80% and test 20%
X_train, X_test, y_train, y_test = train_test_split(features,
                                                    LOS,
                                                    test_size = .20,
                                                    random_state = 42)

xgb_model = XGBRegressor(objective='reg:squarederror', random_state=42)

# Define the parameter grid for grid search
param_grid = {
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'n_estimators': [100, 200, 300],
    'alpha': [0, 0.1, 1],
    'lambda': [0, 0.1, 1]
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5, n

# Perform grid search on the training data
grid_search.fit(X_train, y_train, early_stopping_rounds=10, eval_set=[(X_test,

# Get the best estimator from grid search
best_xgb_model = grid_search.best_estimator_

# Evaluate the best model on the test set
y_pred = best_xgb_model.predict(X_test)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)

# Print the best parameters found
print("Best Parameters:", grid_search.best_params_)

```

C:\Users\User\anaconda3\envs\fyp\lib\site-packages\xgboost\sklearn.py:835: UserWarning: `early_stopping_rounds` in `fit` method is deprecated for better compatibility with scikit-learn, use `early_stopping_rounds` in constructor or `set_params` instead.

warnings.warn(

Best Parameters: {'alpha': 0, 'lambda': 1, 'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 300}

```
In [13]: #fit the model using the best params
refined_xgb_model = XGBRegressor(reg_alpha=0, reg_lambda=1, learning_rate=0.1,
                                max_depth=7, n_estimators=300)

# Fit the model to the training data
refined_xgb_model.fit(X_train, y_train)

y_test_preds = refined_xgb_model.predict(X_test)
r2_optimized = r2_score(y_test, y_test_preds)
mse_optimized = mean_squared_error(y_test, y_test_preds)
mae_optimized = mean_absolute_error(y_test, y_test_preds)
rmse_optimized = mean_squared_error(y_test, y_test_preds, squared=False)
print("Optimized R2 score is: {:.2f}".format(r2_optimized))
print("Optimized MSE score is: {:.2f}".format(mse_optimized))
print("Optimized MAE score is: {:.2f}".format(mae_optimized))
print("Optimized RMSE score is: {:.2f}".format(rmse_optimized))

Optimized R2 score is: 0.639121
Optimized MSE score is: 7.902774
Optimized MAE score is: 1.420043
Optimized RMSE score is: 2.811187
```

```
In [28]: print('Model refinement improved R2 score by {:.4f}'.format(r2_optimized-r2_no
Model refinement improved R2 score by 0.0348
```

Result

```
In [29]: LOS_predict = y_test_preds[:20]
LOS_actual = y_test[:20]
```



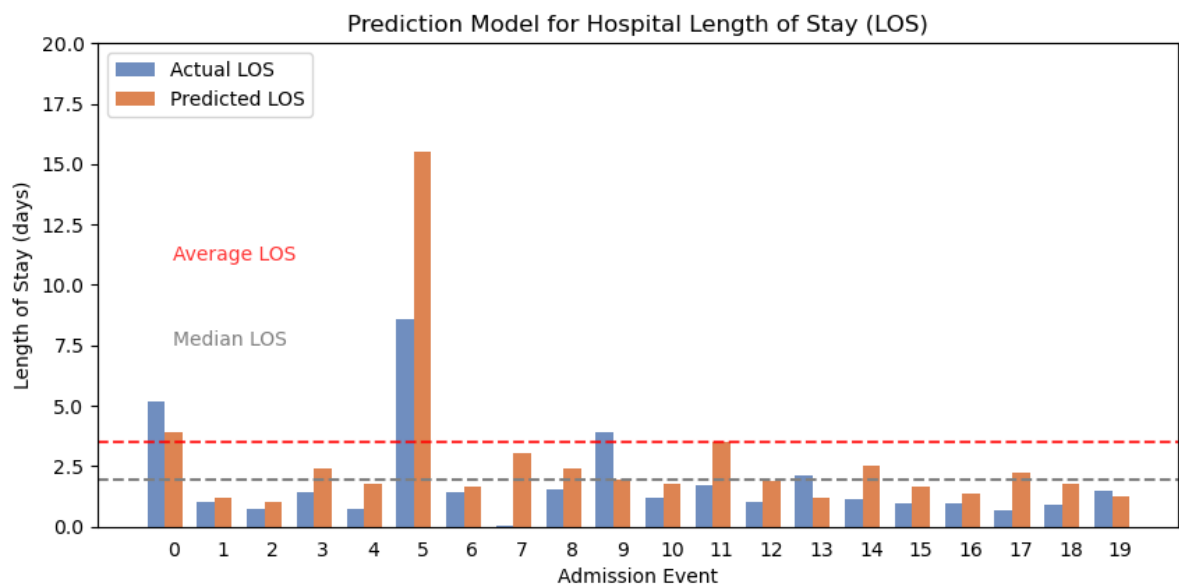
```

In [30]: fig, ax = plt.subplots(figsize=(10, 4.5))

ind = np.arange(0,20)
pad = 0.15
width = 0.35
set_actual = ax.bar(pad+ind, LOS_actual, width, color='#4c72b0', alpha=0.8)
set_predict = ax.bar(pad+ind+width, LOS_predict, width, color='#dd8452')

ax.set_ylabel('Length of Stay (days)')
ax.set_xlabel('Admission Event')
ax.set_title('Prediction Model for Hospital Length of Stay (LOS)')
ax.text(0.5, 11, 'Average LOS', fontdict=None, color='red', alpha=0.8)
ax.text(0.5, 7.5, 'Median LOS', fontdict=None, color='gray')
ax.set_xticks(pad + ind + width)
ax.set_ylim(0, 20)
ax.set_xticklabels(list(range(20)))
ax.axhline(y=actual_median_los, xmin=0, xmax=20, ls='--', color='gray')
ax.axhline(y=actual_mean_los, xmin=0, xmax=20, ls='--', color='red', alpha=0.8)
ax.legend( (set_actual, set_predict), ('Actual LOS', 'Predicted LOS'),
          loc='upper left')
ax.tick_params(bottom=False, top=False, right=False)

```



```

In [31]: #save the model
import pickle

filename = './data/model/los_model.sav'
with open(filename, 'wb') as file:
    pickle.dump(refined_xgb_model, file)

```

```

In [ ]:

```